



selected protected attribute - gender
privileged group is assumed to be male

Basic model



The statistical parity and predictive equality for the basic model are not good - they go well into the ranges of $(\epsilon, 1/\epsilon)$ and the equal opportunity statistic is almost in the range, which is good for the model.

another model - no gender included in the data



The fairness statistics are almost the same as in the original model which shows that the potential gender bias is hidden, and not only explicitly defined by the input data.

bias mitigation - resampling

